

UNCLASSIFIED//FOUO



# *Assembly Language* *Linux 86 Assembly* *part 1*



UNCLASSIFIED



# 1. What is Assembly Language?



UNCLASSIFIED



# What is Assembly Language?

- Low--level programming language
- Communicate with microprocessor
- Specific to the processor family
- An almost one--one correspondence with machine code



UNCLASSIFIED



# I only speak binary!



010101010111101011010101010101  
111010101101011010101010101011010  
0101010111101000011110101010101

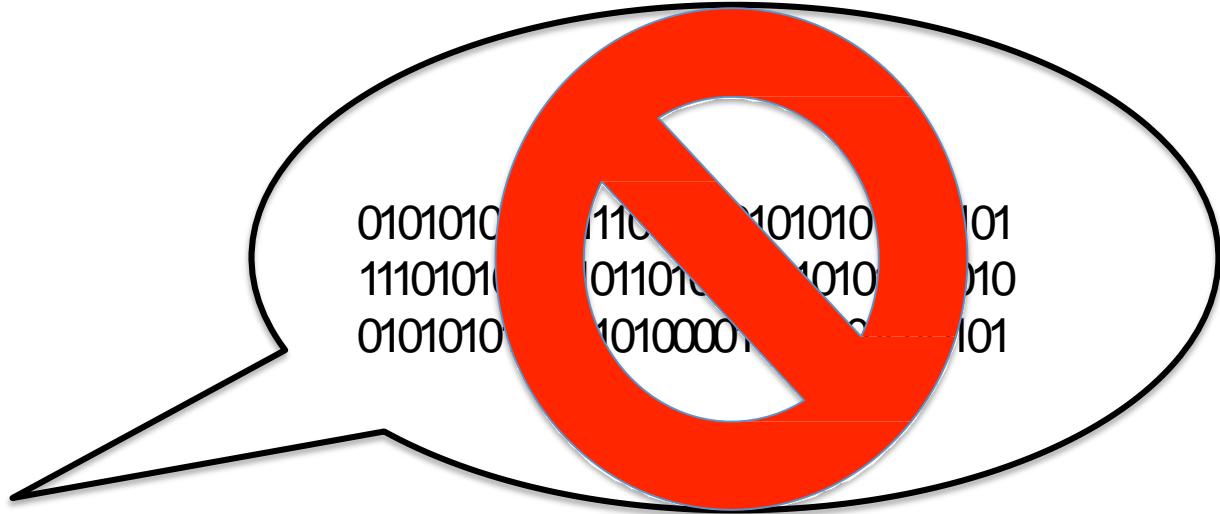
UNCLASSIFIED



**UNCLASSIFIED**



# *Humans cannot speak binary*



**UNCLASSIFIED**



UNCLASSIFIED



# Assembly Language

```
mov eax, ebx  
xor eax, eax  
add eax, 0xff
```

Assembler + Linker

Translator

```
010110100101  
111010101010  
101010101010
```



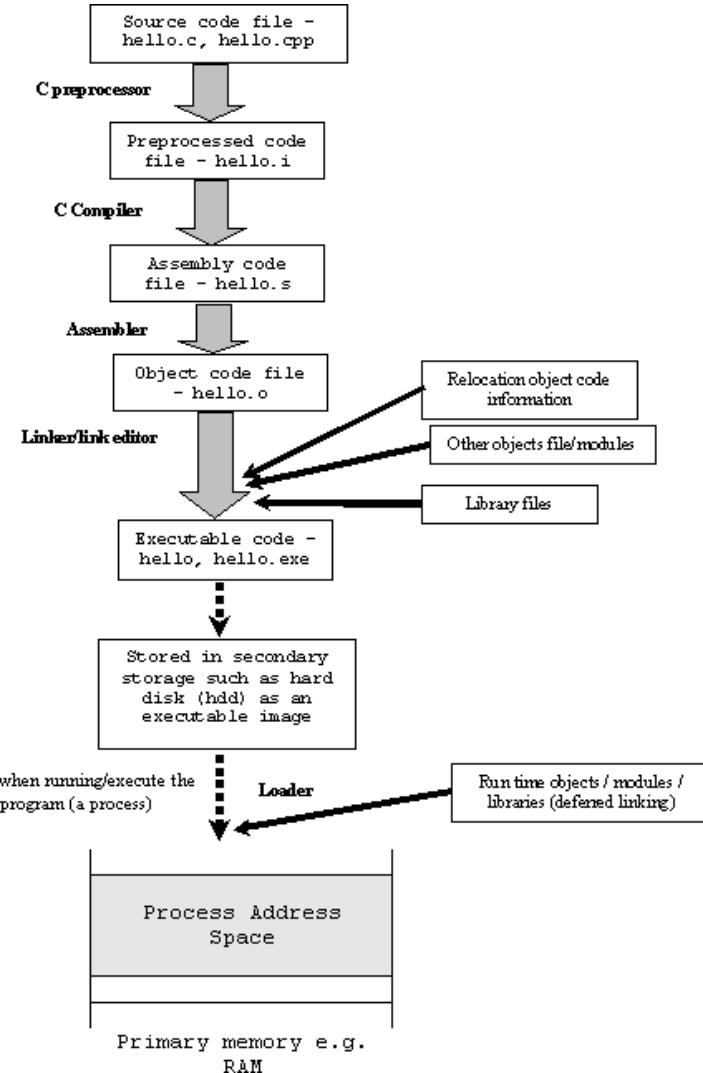
UNCLASSIFIED



UNCLASSIFIED



# Correlation with HLLs



UNCLASSIFIED



UNCLASSIFIED



## *Different Processors – Different Assembly Language*

- Intel
- ARM
- MIPS



UNCLASSIFIED



# *Intel Architecture*

- IA--32
- IA--64



# SecurityTube Linux Assembly Expert<sup>32</sup>

- IA-32 Assembly on the Linux OS



UNCLASSIFIED



# Why IA--32?

- Large number of machines out there still running IA--32
- Logical progression to IA--64
- Shellcoding, Encoders, Decoders, Packers etc. implementation difference



# Exercise 1.2: Understanding your CPU

- Find CPU details on the Ubuntu System
- How do you know if you are on a 32/64 bit CPU?
- How do you know your CPUs additional capabilities such as FPU, MMX, SSE, SSE2 etc.



UNCLASSIFIED



# Module 1: 32--Bit ASM on Linux

## Exercise 1.1

**Topic: What is Assembly Language?**



UNCLASSIFIED



# Exercise 1.1: Lab Setup

- Ubuntu 12.04 LTS 32--bit Edition
- Installed in Virtualbox\vmware



UNCLASSIFIED



# Ubuntu

- 32-bit Ubuntu Desktop Edition 12.10 used for the course

<http://www.ubuntu.com/download/desktop>

- Install Virtualbox

<https://www.virtualbox.org>

/

UNCLASSIFIED



UNCLASSIFIED

# *Installation*

---



- Nasm



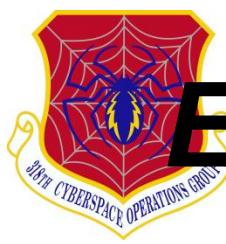
UNCLASSIFIED



# *Module 1: 32--Bit ASM on Linux*

## Exercise 1.2

### Topic: What is Assembly Language?



# Exercise 1.2: Understanding your CPU

- Find CPU details on the Ubuntu System
- How do you know if you are on a 32/64 bit CPU?
- How do you know your CPUs additional capabilities such as FPU, MMX, SSE, SSE2 etc.



UNCLASSIFIED



# Find CPU Details

```
securitytube@securitytube-VirtualBox:~$ lscpu
Architecture:          i686
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                1
On-line CPU(s) list:   0
Thread(s) per core:    1
Core(s) per socket:    1
Socket(s):             1
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 37
Stepping:               5
CPU MHz:               2534.821
BogoMIPS:              5069.64
L1d cache:              32K
L1d cache:              32K
L2d cache:              6144K
securitytube@securitytube-VirtualBox:~$ █
```

UNCLASSIFIED



UNCLASSIFIED



# /proc/cpuinfo

```
securitytube@securitytube-VirtualBox:~$ cat /proc/cpuinfo
processor       : 0
vendor_id      : GenuineIntel
cpu family     : 6
model          : 37
model name     : Intel(R) Core(TM) i5 CPU      M 540 @ 2.53GHz
stepping        : 5
cpu MHz        : 2534.821
cache size     : 6144 KB
fdt_bug        : no
hlt_bug        : no
f00f_bug       : no
coma_bug       : no
fpu             : yes
fpu_exception   : yes
cpuid level    : 5
wp              : yes
flags           : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx
fxsr sse sse2 syscall nx rdtscp lm constant_tsc up pni monitor ssse3 lahf_lm
bogomips        : 5069.64
clflush size    : 64
cache_alignment : 64
address sizes   : 36 bits physical, 48 bits virtual
power management:

securitytube@securitytube-VirtualBox:~$
```

UNCLASSIFIED



UNCLASSIFIED



# *Module 1: 32--Bit ASM on Linux*

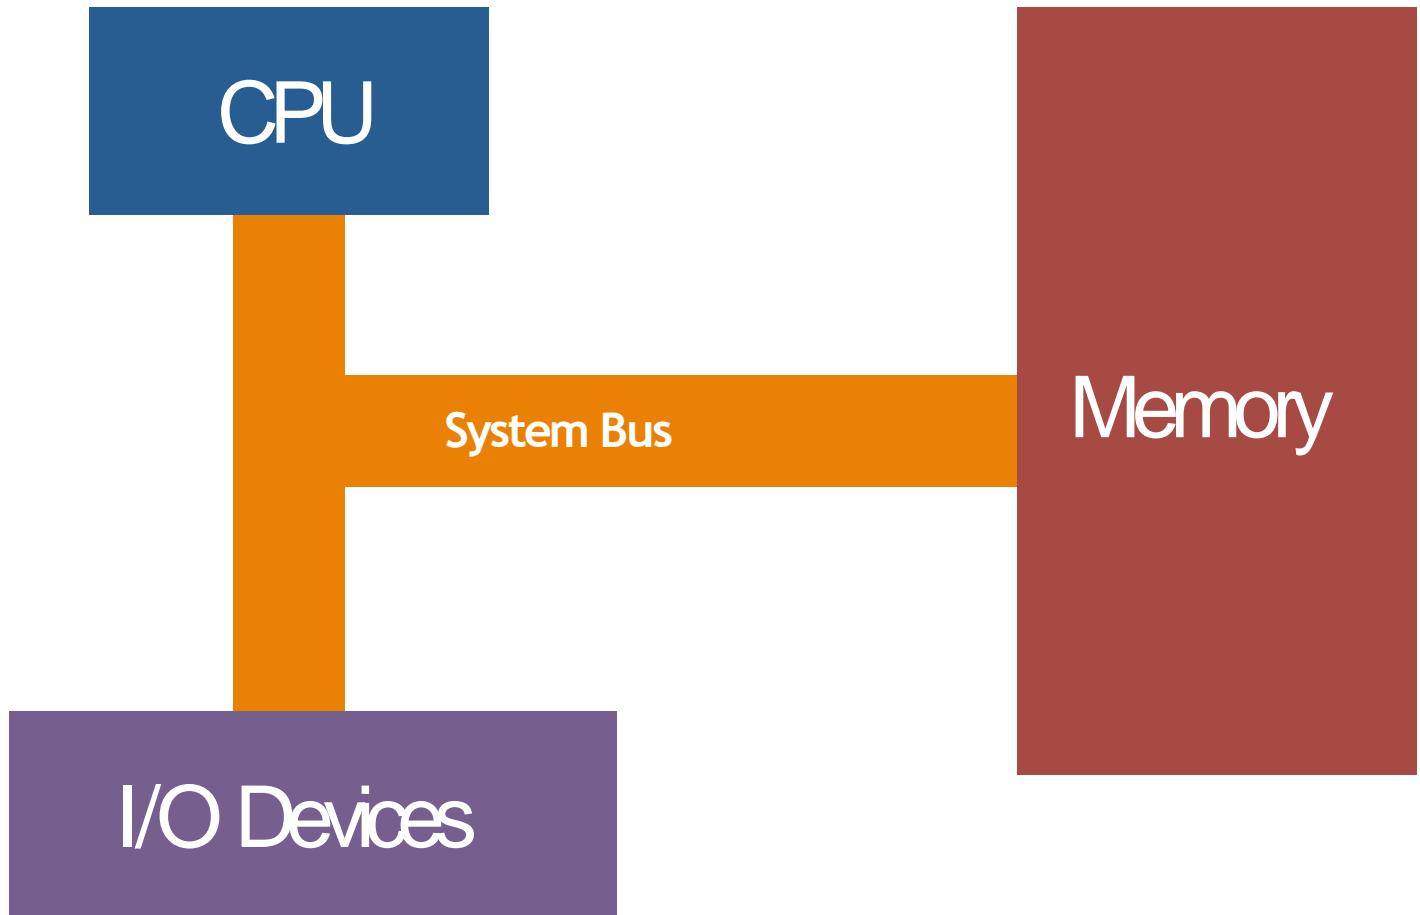
## **2. IA--32 Architecture**



UNCLASSIFIED



# System Organization Basics



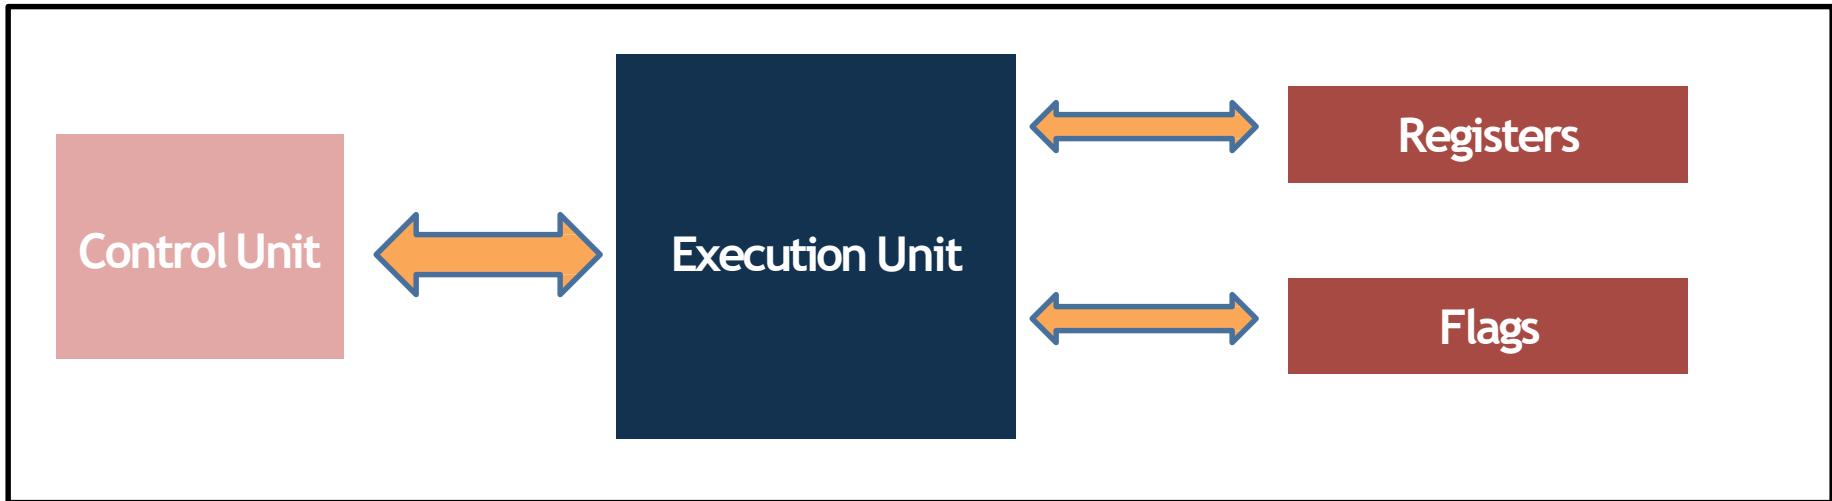
UNCLASSIFIED



UNCLASSIFIED



# CPU



UNCLASSIFIED



UNCLASSIFIED



# IA--32 Registers (Logical Diagram)

General Purpose  
Registers

Segment Registers

Flags, EIP

Floating Point Unit  
Registers

MMX  
Registers

XMM  
Registers

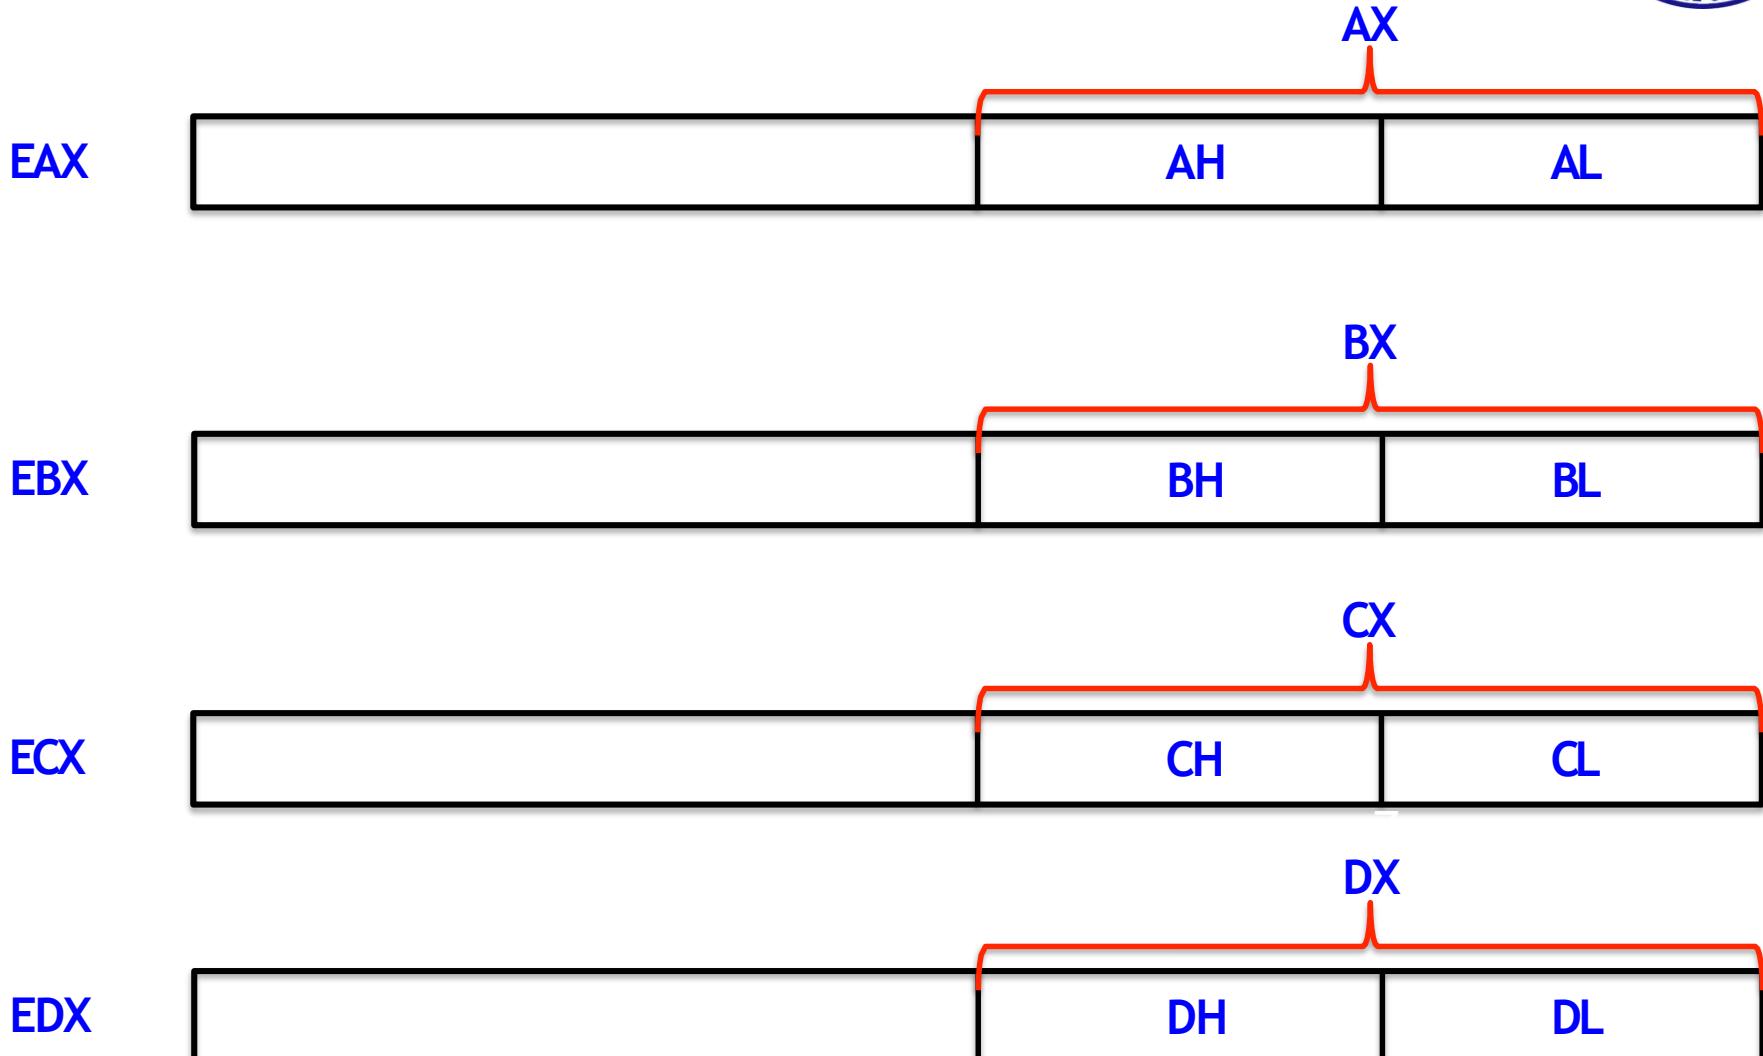
UNCLASSIFIED



UNCLASSIFIED



# General Purpose Registers



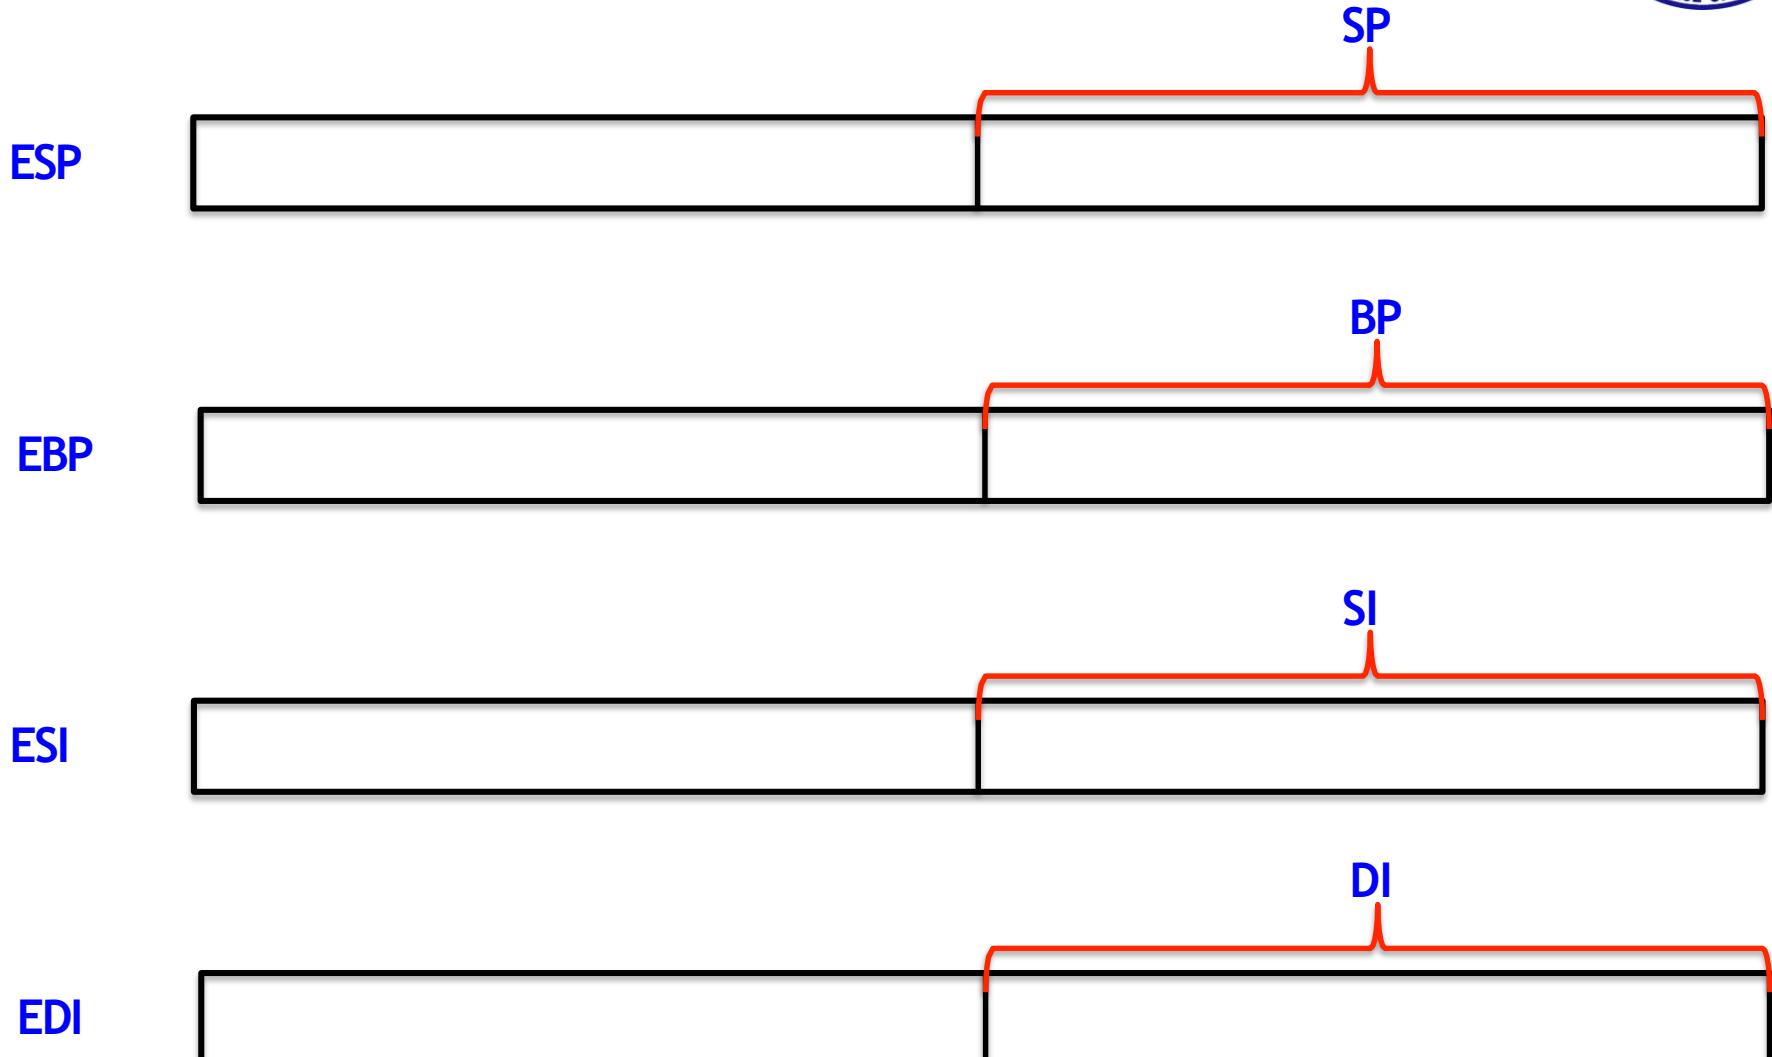
UNCLASSIFIED



UNCLASSIFIED



# General Purpose Registers



UNCLASSIFIED



UNCLASSIFIED

# CPA Common Functionality



EAX

Accumulator Register – used for storing operands and result data

EBX

Base Register – Pointer to Data

ECX

Counter Register – Loop operations

EDX

Data Register – I/O Pointer

ESI

EDI      Data Pointer Registers for memory

ESP

operations Stack Pointer Register

EB  
P

Stack Data Pointer  
Register

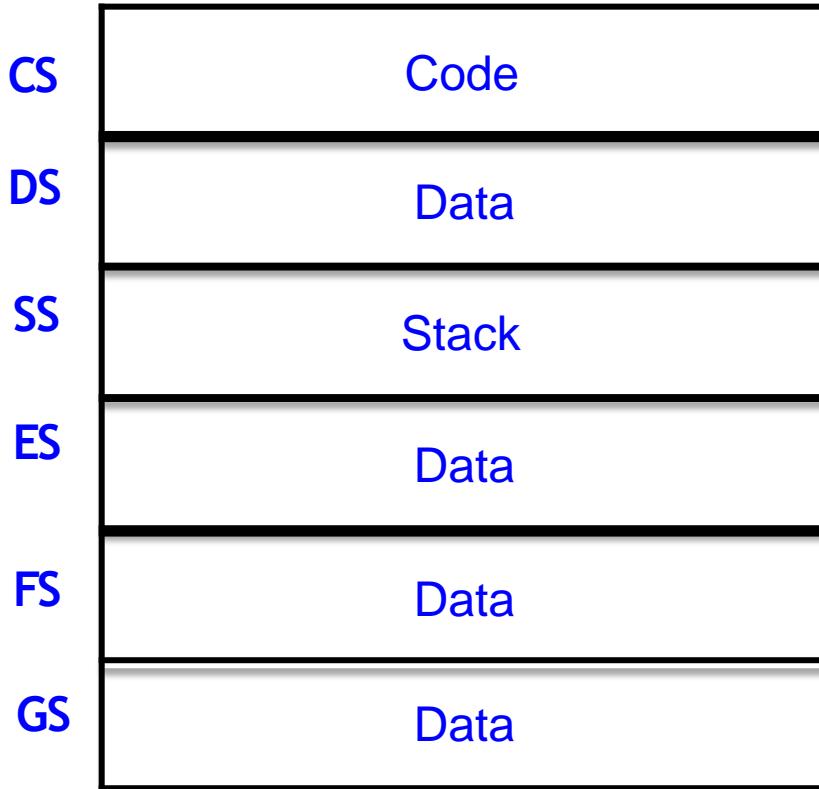
UNCLASSIFIED



UNCLASSIFIED



# Segment Registers



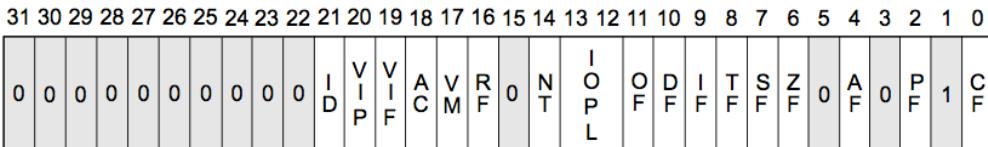
UNCLASSIFIED



UNCLASSIFIED



# EFLAGS Register



- X ID Flag (ID) \_\_\_\_\_
- X Virtual Interrupt Pending (VIP) \_\_\_\_\_
- X Virtual Interrupt Flag (VIF) \_\_\_\_\_
- X Alignment Check (AC) \_\_\_\_\_
- X Virtual-8086 Mode (VM) \_\_\_\_\_
- X Resume Flag (RF) \_\_\_\_\_
- X Nested Task (NT) \_\_\_\_\_
- X I/O Privilege Level (IOPL) \_\_\_\_\_
- S Overflow Flag (OF) \_\_\_\_\_
- C Direction Flag (DF) \_\_\_\_\_
- X Interrupt Enable Flag (IF) \_\_\_\_\_
- X Trap Flag (TF) \_\_\_\_\_
- S Sign Flag (SF) \_\_\_\_\_
- S Zero Flag (ZF) \_\_\_\_\_
- S Auxiliary Carry Flag (AF) \_\_\_\_\_
- S Parity Flag (PF) \_\_\_\_\_
- S Carry Flag (CF) \_\_\_\_\_

S Indicates a Status Flag

C Indicates a Control Flag

X Indicates a System Flag

Reserved bit positions. DO NOT USE.  
Always set to values previously read.

UNCLASSIFIED



UNCLASSIFIED

EIP



EIP

31

- Instruction Pointer
- Holy grail for Shellcoding, Exploit Research etc.



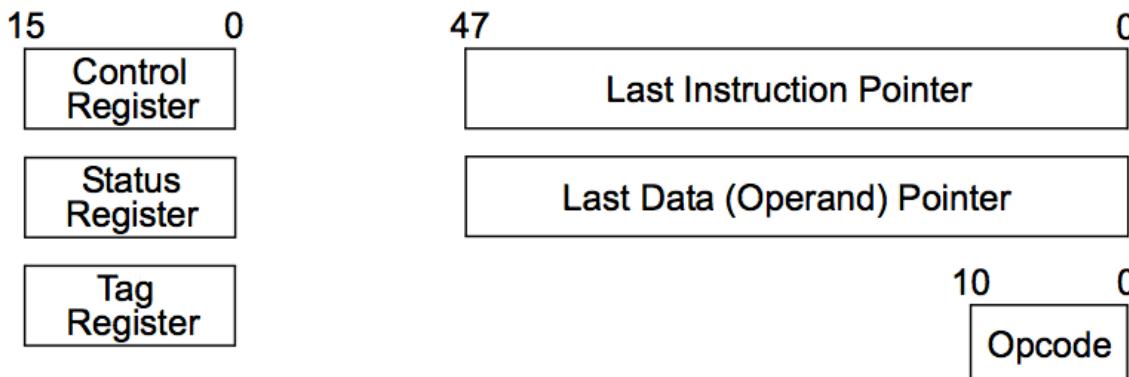
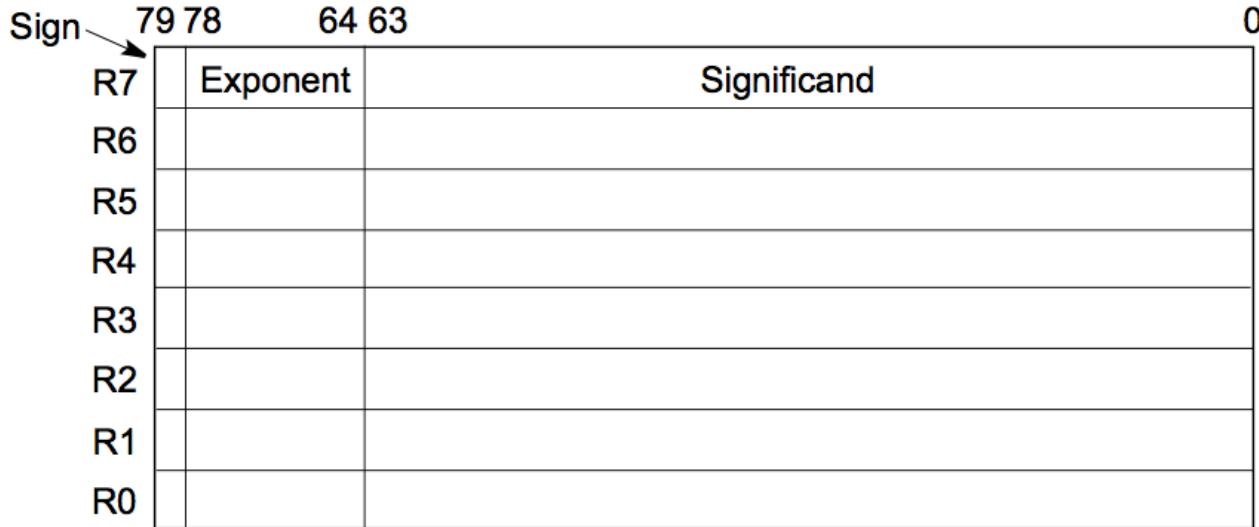
UNCLASSIFIED



# Floating Point Unit (FPU) or x87

## Data Registers

ST(0)to ST(7)



Reference: Intel Manual

UNCLASSIFIED



UNCLASSIFIED

# SIMD



- Single Instruction Multiple Data
- Extensions
  - MMX
  - SSE
  - SSE2
  - SSE3
- Uses MMX and XMM Registers



UNCLASSIFIED

**MMX**

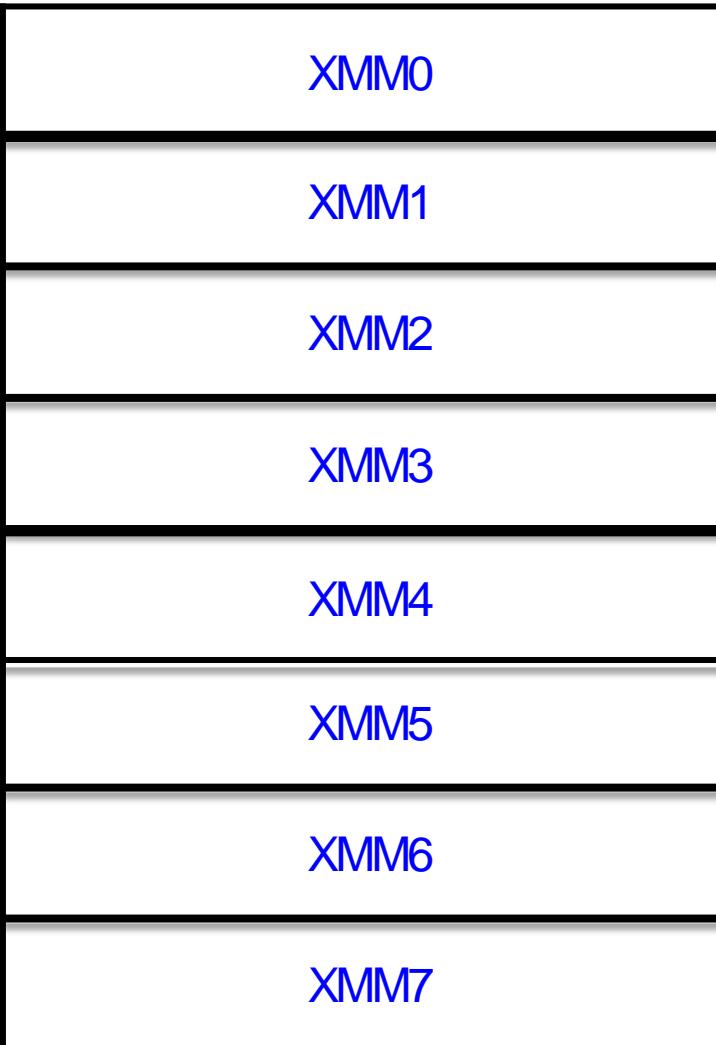


ST(0)	MM0
ST(1)	MM1
ST(2)	MM2
ST(3)	MM3
ST(4)	MM4
ST(5)	MM5
ST(6)	MM6
ST(7)	MM7

UNCLASSIFIED



UNCLASSIFIED



UNCLASSIFIED



UNCLASSIFIED



# Exercise 12.1: Lab Setup

- Inspect the General Purpose, Segment, EFLAGS, FPU, MMX, XMM etc. registers on your Ubuntu system



UNCLASSIFIED



# Exercise 12.1: Lab Setup

- Inspect the General Purpose, Segment, EFLAGS, FPU, MMX, XMM etc. registers on your Ubuntu system



UNCLASSIFIED



# *Module 1: 32--Bit ASM on Linux*

## 3. CPU Modes and Memory Management



# CPU Modes for IA--32

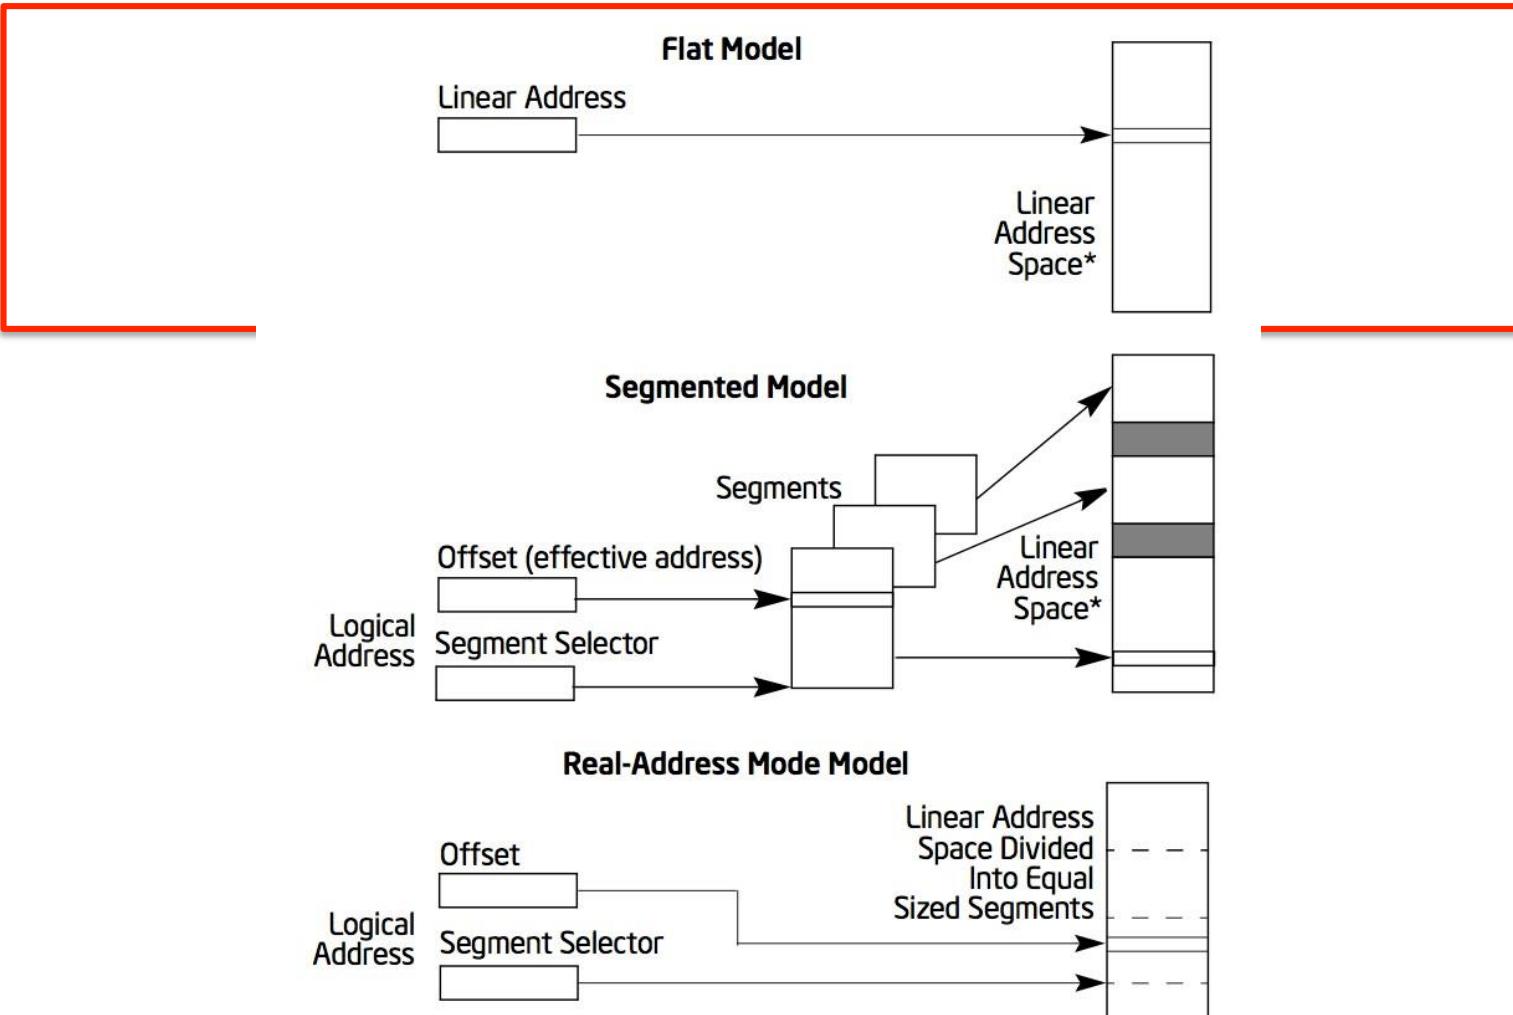
- Real Mode
  - At power up or reset
  - Can only access 1 MB memory
  - No memory protection
  - Privilege Levels (Kernel vs User space) not possible
- Protected Mode
  - Up to 4GB memory
  - memory protection / privilege level / multi-tasking
  - Supports Virtual-8086 mode
- System Management Mode
  - Used for power management tasks



UNCLASSIFIED



# Memory Models



Source: Intel manuals

UNCLASSIFIED



UNCLASSIFIED



# Linux: CPU Mode and Memory Model

- 32--Bit Linux uses:
  - Protected Mode
  - Flat Memory model
  - 4GB Addressable space =>  $2^{32}$
  - Memory Protection
  - Privilege Levels of Code
  - Segment registers point to segment descriptors
    - GDT / LDT / IDT (Global / Local / Interrupt )

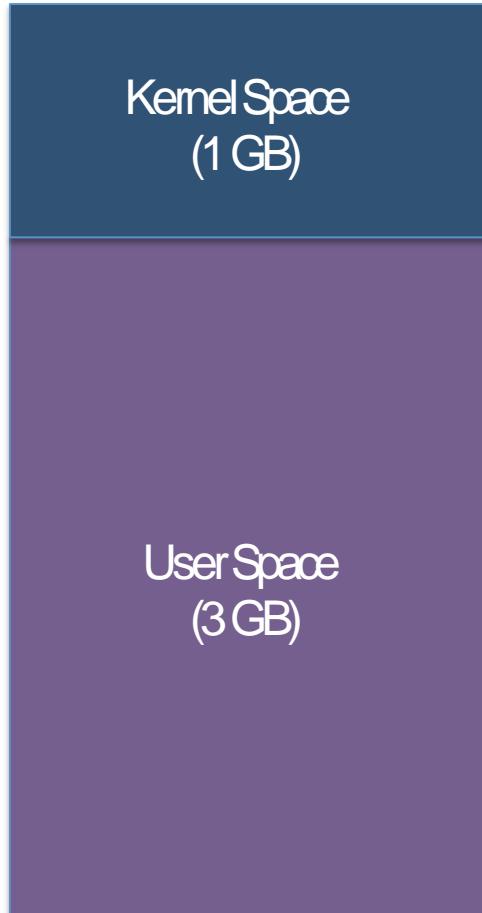
UNCLASSIFIED



UNCLASSIFIED



# *Virtual Memory Model*



UNCLASSIFIED



UNCLASSIFIED

0xFFFFF

FFF

0xC00000

00

0x0804800

0

0x00000000

0



Function Args + Local  
Vars

Dynamic  
Memory

Uninitialized  
Data

Initialized  
Data

Program  
Code



UNCLASSIFIED



UNCLASSIFIED



# *View Process Organization*

- /Proc
  - /proc/pid/maps
- pmap
- Attach and view using GDB



UNCLASSIFIED



# cat /proc/pid/maps

```
securitytube@securitytube-VirtualBox:~$ cat /proc/self/maps
08048000-08053000 r-xp 00000000 08:01 262170      /bin/cat
08053000-08054000 r--p 0000a000 08:01 262170      /bin/cat
08054000-08055000 rw-p 0000b000 08:01 262170      /bin/cat
0917e000-0919f000 rw-p 00000000 00:00 0          [heap]
b73cd000-b75cd000 r--p 00000000 08:01 140017      /usr/lib/locale/locale-archive
b75cd000-b75ce000 rw-p 00000000 00:00 0
b75ce000-b776d000 r-xp 00000000 08:01 265843      /lib/i386-linux-gnu/libc-2.15.so
b776d000-b776f000 r--p 0019f000 08:01 265843      /lib/i386-linux-gnu/libc-2.15.so
b776f000-b7770000 rw-p 001a1000 08:01 265843      /lib/i386-linux-gnu/libc-2.15.so
b7770000-b7773000 rw-p 00000000 00:00 0
b7782000-b7783000 r--p 005db000 08:01 140017      /usr/lib/locale/locale-archive
b7783000-b7785000 rw-p 00000000 00:00 0
b7785000-b7786000 r-xp 00000000 00:00 0          [vds]
b7786000-b77a6000 r-xp 00000000 08:01 265823      /lib/i386-linux-gnu/ld-2.15.so
b77a6000-b77a7000 r--p 0001f000 08:01 265823      /lib/i386-linux-gnu/ld-2.15.so
b77a7000-b77a8000 rw-p 00020000 08:01 265823      /lib/i386-linux-gnu/ld-2.15.so
bfdfc000-bfe1d000 rw-p 00000000 00:00 0          [stack]
securitytube@securitytube-VirtualBox:~$ █
```

UNCLASSIFIED

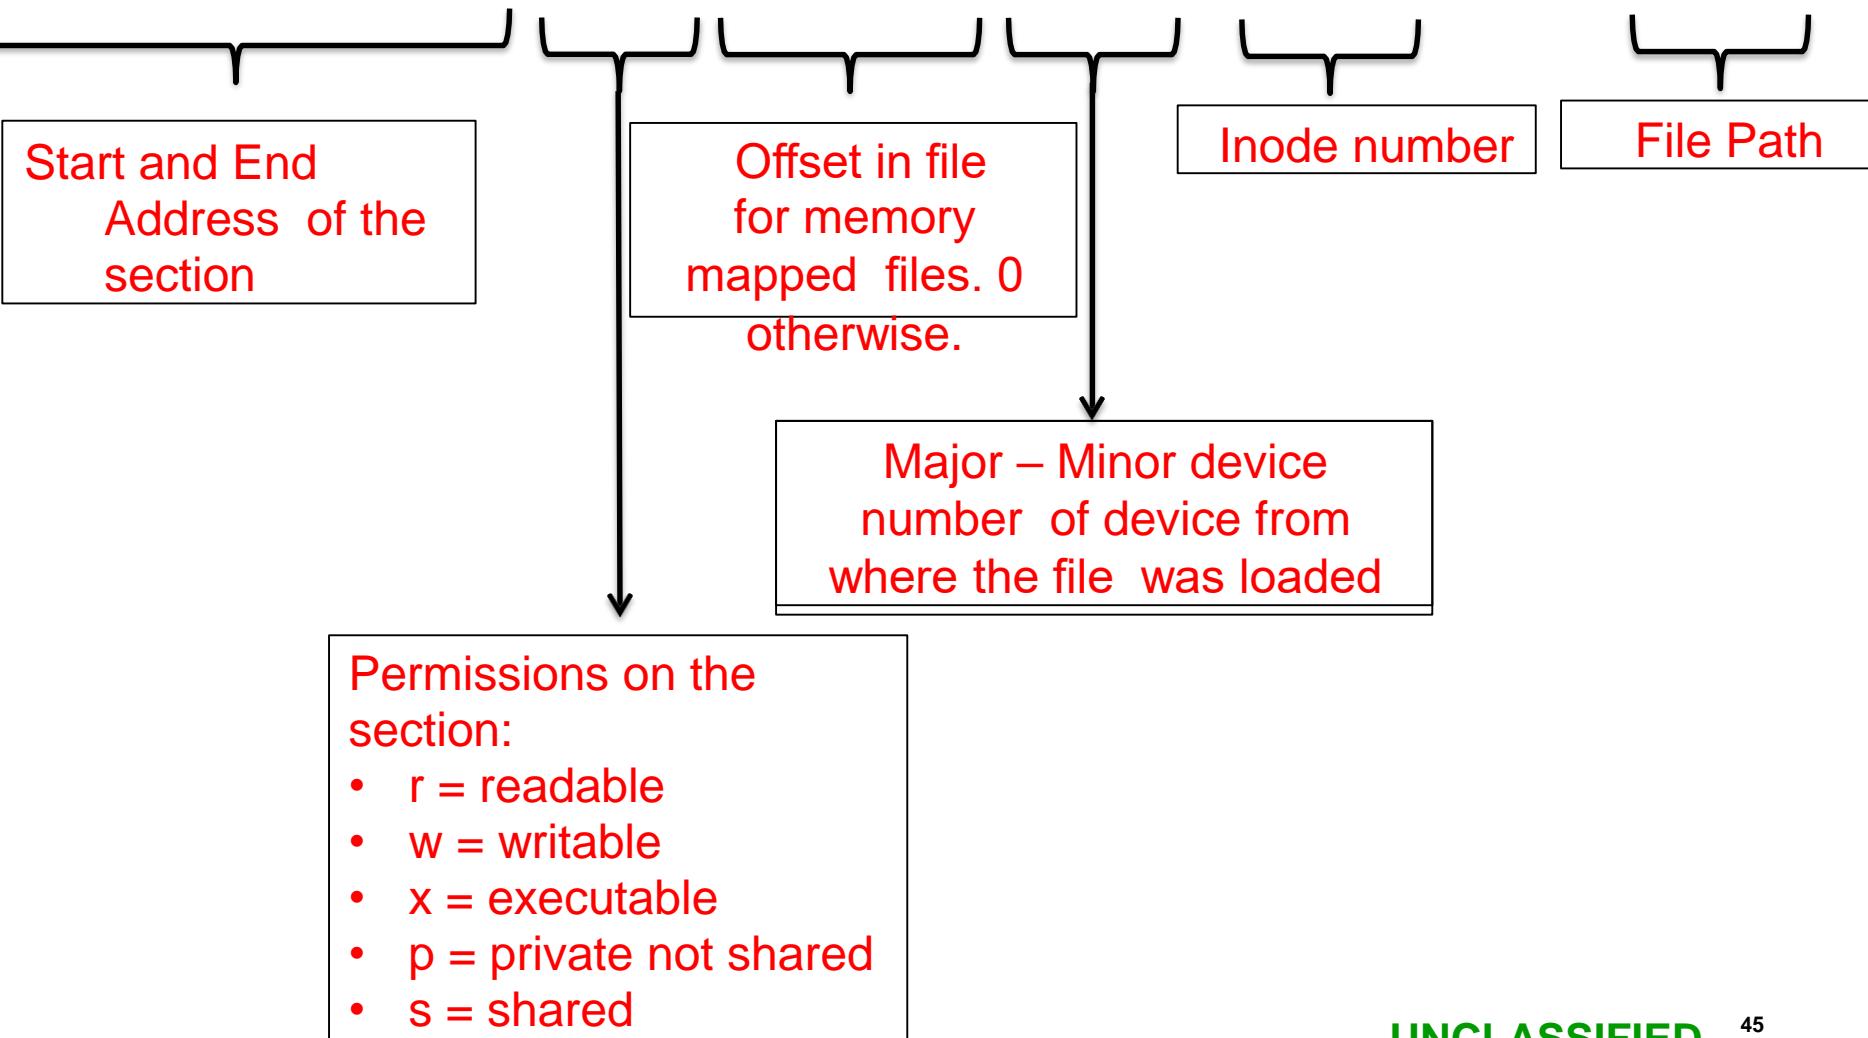


UNCLASSIFIED



# What does all this mean?

08048000-08053000 r-xp 00000000 08:01 262170 /bin/cat



UNCLASSIFIED



UNCLASSIFIED



# Process Map within GDB

(gdb) info proc mappings

process 2837

Mapped address spaces:

Start Addr	End Addr	Size	Offset	objfile
0x8048000	0x8124000	0xdc000	0x0	/bin/bash
0x8124000	0x8125000	0x1000	0xdb000	/bin/bash
0x8125000	0x812a000	0x5000	0xdc000	/bin/bash
0x812a000	0x812f000	0x5000	0x0	[heap]
0xb7e00000	0xb7e01000	0x1000	0x0	
0xb7e01000	0xb7fa0000	0x19f000	0x0	/lib/i386-linux-gnu/libc-2.15.so
0xb7fa0000	0xb7fa2000	0x2000	0x19f000	/lib/i386-linux-gnu/libc-2.15.so
0xb7fa2000	0xb7fa3000	0x1000	0x1a1000	/lib/i386-linux-gnu/libc-2.15.so
0xb7fa3000	0xb7fa7000	0x4000	0x0	
0xb7fa7000	0xb7faa000	0x3000	0x0	/lib/i386-linux-gnu/libdl-2.15.so
0xb7faa000	0xb7fab000	0x1000	0x2000	/lib/i386-linux-gnu/libdl-2.15.so
0xb7fab000	0xb7fac000	0x1000	0x3000	/lib/i386-linux-gnu/libdl-2.15.so
0xb7fac000	0xb7fc8000	0x1c000	0x0	/lib/i386-linux-gnu/libtinfo.so.5.9
0xb7fc8000	0xb7fca000	0x2000	0x1b000	/lib/i386-linux-gnu/libtinfo.so.5.9
0xb7fca000	0xb7fcbb000	0x1000	0x1d000	/lib/i386-linux-gnu/libtinfo.so.5.9
0xb7fdb000	0xb7fdd000	0x2000	0x0	
0xb7fdd000	0xb7fde000	0x1000	0x0	[vds]
0xb7fde000	0xb7ffe000	0x20000	0x0	/lib/i386-linux-gnu/ld-2.15.so
0xb7ffe000	0xb7fff000	0x1000	0x1f000	/lib/i386-linux-gnu/ld-2.15.so
0xb7fff000	0xb8000000	0x1000	0x20000	/lib/i386-linux-gnu/ld-2.15.so
0xbffdf000	0xc0000000	0x21000	0x0	[stack]

(gdb)

UNCLASSIFIED



UNCLASSIFIED



# *Module 1: 32--Bit ASM on Linux*

## 4. Hello World



UNCLASSIFIED



# IA--32 Instruction Set

- General Purpose Instructions
- x87 FPU Instructions
- MMX / SSE / SSE2/ SSE3 / SSE4 Instructions
- Other Instruction Set extensions



UNCLASSIFIED



# Programming in Assembly

- NASM + LD for assembling and linking
- Executable in ELF format

NASM Documentation:

<http://nasm.us/>

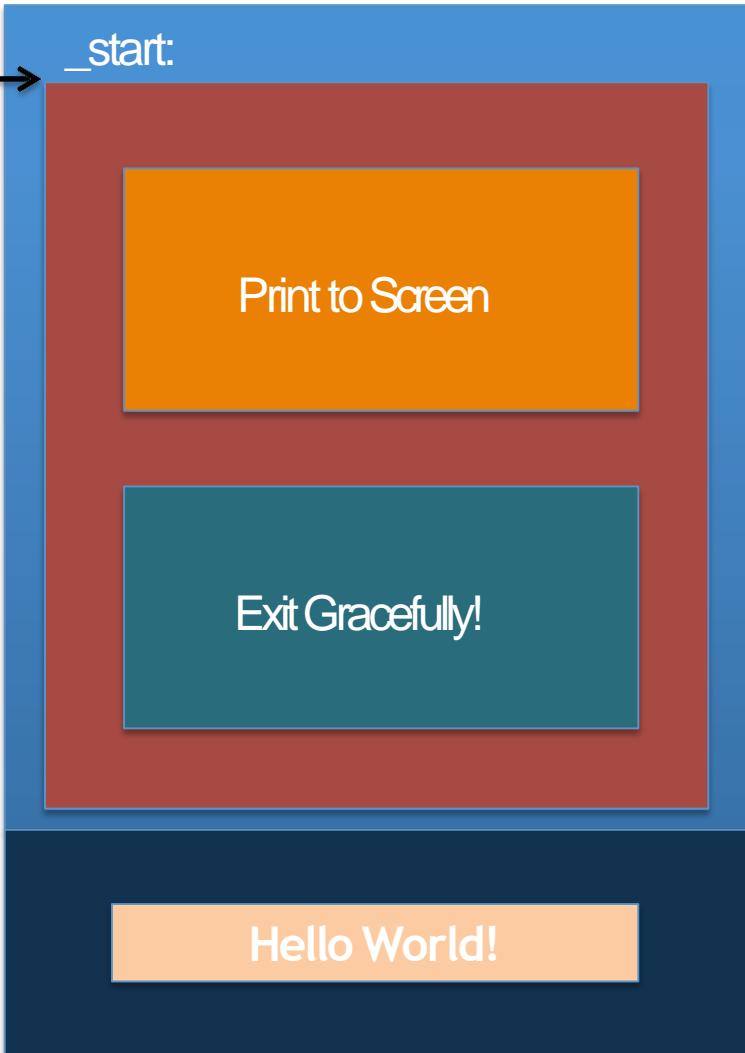


UNCLASSIFIED



# Hello World!

Entry Point of  
Program?



UNCLASSIFIED



# Why System Calls?

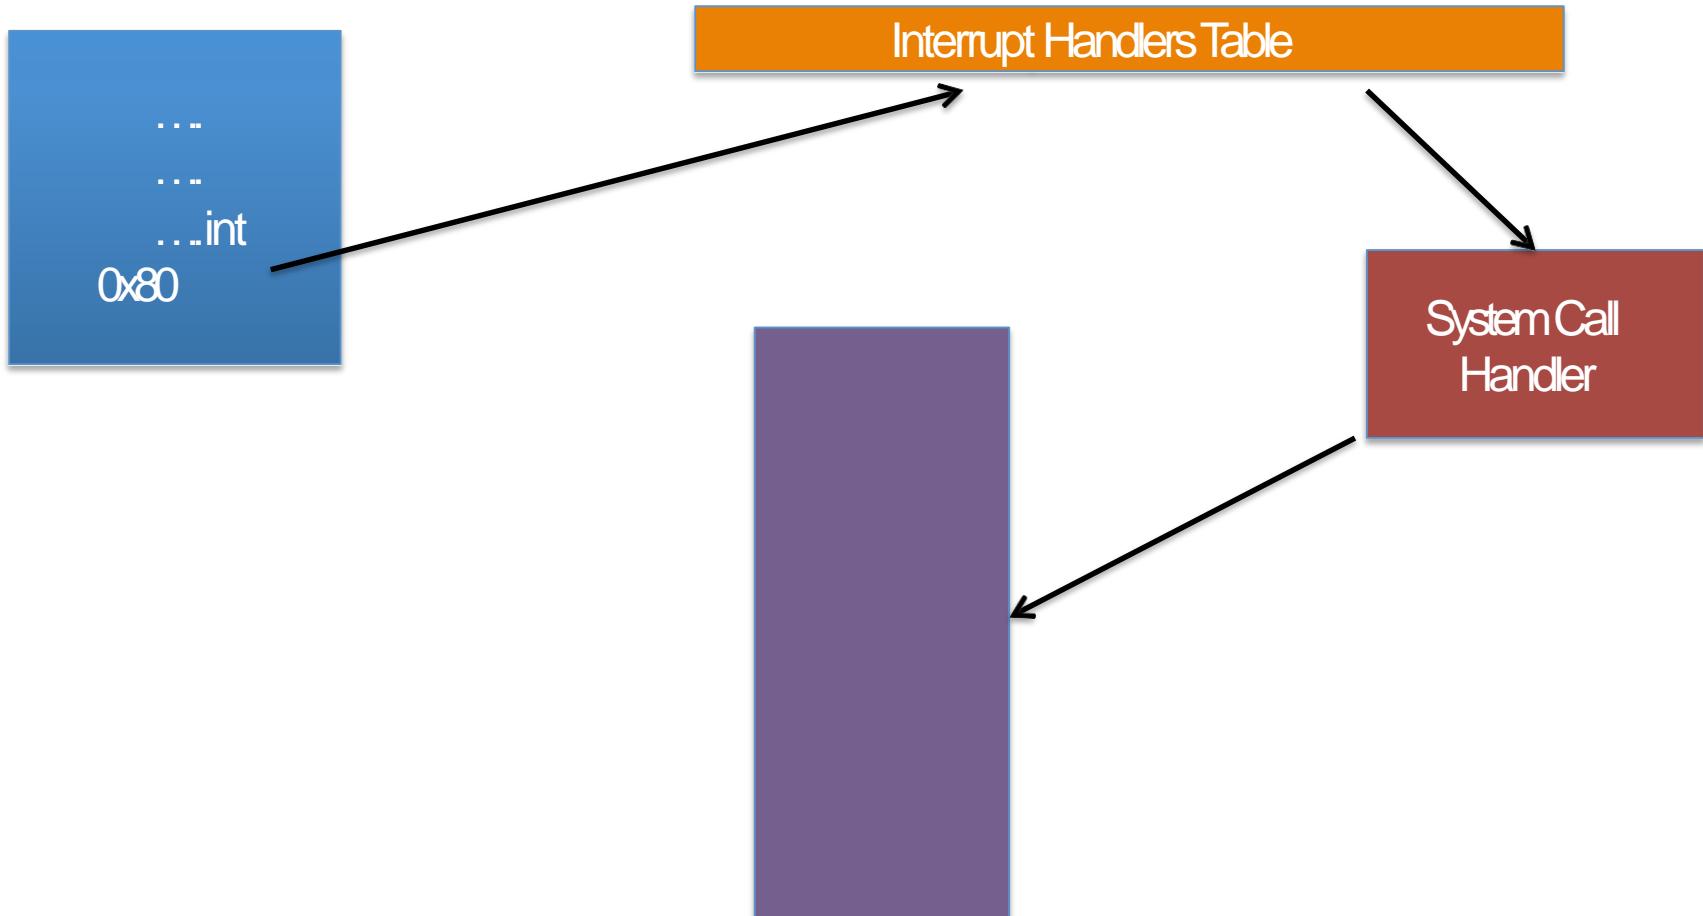
- Leverage OS for tasks
- Imagine if you had to write code from scratch to:
  - write to disk
  - print on screen
  - ...
- System Calls provide a simple interface for user space programs to the Kernel



UNCLASSIFIED



# How do System calls work?



UNCLASSIFIED



UNCLASSIFIED



# IA--32 Mechanism to invoke System Call

- int 0x80
- SYSENTER

Modern implementations using  
VDSO [Virtual Dynamic Shared  
Object]

[http://articles.manugarg.com/systemcallinlinux2\\_6.html](http://articles.manugarg.com/systemcallinlinux2_6.html)

UNCLASSIFIED



# Where are these system calls defined?

```
#define __NR_restart_syscall          0
#define __NR_exit                      1
#define __NR_fork                      2
#define __NR_read                       3
#define __NR_write                     4
#define __NR_open                       5
#define __NR_close                     6
#define __NR_waitpid                   7
#define __NR_creat                     8
#define __NR_link                      9
#define __NR_unlink                    10
#define __NR_execve                    11
#define __NR_chdir                     12
#define __NR_time                      13
#define __NR_mknod                     14
#define __NR_chmod                     15
#define __NR_lchown                    16
#define __NR_break                     17
#define __NR_oldstat                   18
#define __NR_lseek                     19
#define __NR_getpid                   20
#define __NR_mount                     21
#define __NR_umount                   22
#define __NR_setuid                   23
#define __NR_getuid                   24
#define __NR_stime                     25
"/usr/include/i386-linux-gnu/asm/unistd_32.h"
```



UNCLASSIFIED

# write()



WRITE(2)

Linux Programmer's Manual

## NAME

**write** - write to a file descriptor

## SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

## DESCRIPTION

**write()** writes up to count bytes from the buffer pointed buf to the

UNCLASSIFIED



UNCLASSIFIED

# exit



\_EXIT(2)

Linux Pro

## NAME

`_exit, _Exit` - terminate the calling process

## SYNOPSIS

```
#include <unistd.h>
```

```
void _exit(int status);
```

UNCLASSIFIED



UNCLASSIFIED



# Invoking System Call with 0x80

EAX	System Call Number	Return Value in EAX
EBX	1st Argument	
ECX	2nd Argument	
EDX	3rd Argument	
ESI	4th Argument	
EDI	5th Argument	

UNCLASSIFIED



UNCLASSIFIED



# Calling Write

WRITE(2)

Linux Programmer's Manual

## NAME

`write` - write to a file descriptor

## SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

## DESCRIPTION

`write()` writes up to `count` bytes from the buffer pointed `buf` to the file descriptor `fd`.

EAX=system call number

ECX=Pointer to "HelloWorld"

EDX=Length of "HelloWorld"

EBX=STDOUT

UNCLASSIFIED



UNCLASSIFIED



# Calling Exit

\_EXIT(2)

Linux Pro

## NAME

\_exit, \_Exit - terminate the calling process

## SYNOPSIS

```
#include <unistd.h>
```

```
void _exit(int status);
```



EAX=System call number

EBX=StatusCode

UNCLASSIFIED



UNCLASSIFIED



## Exercise 1.4.1: GDB

- Use GDB to step through the Hello World program and observe:
  - CPU Registers
  - Memory Location
  - ...



UNCLASSIFIED



# Exercise 1.4.1: GDB

- Use GDB to step through the Hello World program and observe:
  - CPU Registers
  - Memory Location
  - ...